

Q6. Take a regression dataset from Kaggle and implement linear regression.**AIM**

To implement linear regression model on housing data from Kaggle.

SOFTWARE USED

Jupyter Platform - Python Programming Language

PROGRAM CODE

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import mpl_toolkits
%matplotlib inline

data = pd.read_csv("kc_house_data.csv")

data.head()
data.describe()
train1 = data.drop(['id', 'price'],axis=1)

train1.head()

from sklearn.linear_model import LinearRegression

reg = LinearRegression()

labels = data['price']
conv_dates = [1 if values == 2014 else 0 for values in data.date ]
data['date'] = conv_dates
train1 = data.drop(['id', 'price'],axis=1)

from sklearn.model_selection import train_test_split

x_train , x_test , y_train , y_test = train_test_split(train1 , labels , test_size = 0.10,random_state =2)

reg.fit(x_train,y_train)

reg.score(x_test,y_test)
```

OUTPUT

Jupyter Linear_Regression_housesales Last Checkpoint: a minute ago (autosaved) Python 3 (ipykernel)

```
In [35]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import mpl_toolkits
%matplotlib inline
```

```
In [36]: data = pd.read_csv("kc_house_data.csv")
```

```
In [37]: data.head()
```

```
Out[37]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	0	...	7	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	0	...	7	
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	0	...	6	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	0	...	7	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	0	...	8	

5 rows × 21 columns

```
In [38]: data.describe()
```

```
Out[38]:
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
--	----	-------	----------	-----------	-------------	----------	--------	------------

Jupyter Linear_Regression_housesales Last Checkpoint: a minute ago (autosaved) Python 3 (ipykernel)

```
In [38]: data.describe()
```

```
Out[38]:
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
count	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+04	21613.000000	21613.000000
mean	4.580302e+09	5.400881e+05	3.370842	2.114757	2079.899736	1.510697e+04	1.494309	0.007542
std	2.876566e+09	3.671272e+05	0.930062	0.770163	918.440897	4.142051e+04	0.539989	0.086517
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+02	1.000000	0.000000
25%	2.123049e+09	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	1.000000	0.000000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	2.000000	0.000000
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1.000000

```
In [47]: train1 = data.drop(['id', 'price'],axis=1)
```

```
In [48]: train1.head()
```

```
Out[48]:
```

	date	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_ba
0	20141013T000000	3	1.00	1180	5650	1.0	0	0	3	7	1180	
1	20141209T000000	3	2.25	2570	7242	2.0	0	0	3	7	2170	
2	20150225T000000	2	1.00	770	10000	1.0	0	0	3	6	770	
3	20141209T000000	4	3.00	1960	5000	1.0	0	0	5	7	1050	
4	20150218T000000	3	2.00	1680	8080	1.0	0	0	3	8	1680	

```

3 20141209T000000    4    3.00    1960    5000    1.0    0    0    5    7    1050
4 20150218T000000    3    2.00    1680    8080    1.0    0    0    3    8    1680

In [53]: from sklearn.linear_model import LinearRegression

In [54]: reg = LinearRegression()

In [55]: labels = data['price']
conv_dates = [1 if values == 2014 else 0 for values in data.date ]
data['date'] = conv_dates
train1 = data.drop(['id', 'price'],axis=1)

In [56]: from sklearn.model_selection import train_test_split

In [57]: x_train , x_test , y_train , y_test = train_test_split(train1 , labels , test_size = 0.10,random_state

In [58]: reg.fit(x_train,y_train)

Out[58]: LinearRegression()

In [68]: reg.score(x_test,y_test)

Out[68]: 0.7320342760357641

```

DISCUSSION and CONCLUSION

The linear regression model has been applied and executed successfully on housing data.

**Q7. Take a classification dataset from Kaggle and classify the data into output classes.
Also evaluate the classifier efficiency using various evaluation measures.**

AIM

Implement a classification problem on a Kaggle dataset using logistic regression.

SOFTWARE USED

Jupyter Platform - Python Programming Language

PROGRAM CODE

```

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # data visualization
import seaborn as sns # statistical data visualization
%matplotlib inline

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

data = '/kaggle/input/weather-dataset-rattle-package/weatherAUS.csv'

```

```
df = pd.read_csv(data)

# preview the dataset

df.head()

col_names = df.columns

col_names

X = df.drop(['RainTomorrow'], axis=1)

y = df['RainTomorrow']

# split X and y into training and testing sets

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)

# check the shape of X_train and X_test

X_train.shape, X_test.shape
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)

X_train = pd.DataFrame(X_train, columns=[cols])
X_test = pd.DataFrame(X_test, columns=[cols])

X_train.describe()

# train a logistic regression model on the training set
from sklearn.linear_model import LogisticRegression

# instantiate the model
logreg = LogisticRegression(solver='liblinear', random_state=0)

# fit the model
logreg.fit(X_train, y_train)

y_pred_test = logreg.predict(X_test)

y_pred_test

# probability of getting output as 0 - no rain
```

```
logreg.predict_proba(X_test)[:,0]

# probability of getting output as 1 - rain

logreg.predict_proba(X_test)[:,1]

from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred_test)))

y_pred_train = logreg.predict(X_train)

y_pred_train

print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_train)))

# Print the Confusion Matrix and slice it into four pieces

from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred_test)

print('Confusion matrix\n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])


from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred_test))


# print classification accuracy

classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)

print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))

# print classification error

classification_error = (FP + FN) / float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))

# print precision score
```

jupyter logistic-regression-classifier-tutorial Last Checkpoint: 7 minutes ago (autosaved)

[illegible]



Logout

File Edit View Insert Cell Kernel Help

Not Trusted

Python 3 (ipykernel) 








50%	0.479810	0.517958	0.000000	0.220183	0.586207	0.255814	0.236364	0.333333	0.700000	0.520000
75%	0.593824	0.623819	0.187500	0.247706	0.600000	0.310078	0.345455	0.421053	0.830000	0.650000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 118 columns

We now have '`X_train`' dataset ready to be fed into the Logistic Regression classifier. I will do it as follows.

Model training

```
In [99]: # train a Logistic regression model on the training set
from sklearn.linear_model import LogisticRegression

# instantiate the model
logreg = LogisticRegression(solver='liblinear', random_state=0)

# fit the model
logreg.fit(X_train, y_train)


Out[99]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l2',
                             random_state=0, solver='liblinear', tol=0.0001, verbose=0,
                             warm_start=False)
```



Logout

File Edit View Insert Cell Kernel Help

Not Trusted

Python 3 (ipykernel) 








```
random_state=0, solver='liblinear', tol=0.0001, verbose=0,
warm_start=False)
```

Predict results

```
In [100]: y_pred_test = logreg.predict(X_test)

y_pred_test

Out[100]: array(['No', 'No', 'No', ..., 'No', 'No', 'Yes'], dtype=object)
```

predict_proba method

`predict_proba` method gives the probabilities for the target variable(0 and 1) in this case, in array form.

0 is for probability of no rain and 1 is for probability of rain.

```
In [101]: # probability of getting output as 0 - no rain
logreg.predict_proba(X_test)[:,:0]

Out[101]: array([[0.91382428, 0.83565645, 0.82033915, ..., 0.97674285, 0.79855098,
                  0.30734161]])

In [102]: # probability of getting output as 1 - rain
logreg.predict_proba(X_test)[:,:1]

Out[102]: array([[0.08617572, 0.16434355, 0.17966085, ..., 0.02325715, 0.20144902,
                  0.69265839]])
```

```
logreg.predict_proba(X_test)[:,-1]
```

```
Out[102]: array([0.08617572, 0.16434355, 0.17966085, ..., 0.02325715, 0.20144902,
                0.69265839])
```

Check accuracy score

```
In [103]: from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred_test)))

Model accuracy score: 0.8502
```

Here, **y_test** are the true class labels and **y_pred_test** are the predicted class labels in the test-set.

Compare the train-set and test-set accuracy

```
In [104]: y_pred_train = logreg.predict(X_train)

y_pred_train

Out[104]: array(['No', 'No', 'No', ..., 'No', 'No', 'No'], dtype=object)

In [105]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_train)))

Training-set accuracy score: 0.8476
```

Confusion matrix

Training-set accuracy score: 0.8476

Confusion matrix

```
In [113]: # Print the Confusion Matrix and slice it into four pieces

from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred_test)

print('Confusion matrix\n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])
```

Confusion matrix

```
[[20892 1175]
 [ 3086 3286]]
```

True Positives(TP) = 20892

True Negatives(TN) = 3286

False Positives(FP) = 1175

False Negatives(FN) = 3086

File Edit View Insert Cell Kernel Help

Not Trusted

Python 3 (ipykernel) 









Code

```
True Negatives(TN) = 3200
False Positives(FP) = 1175
False Negatives(FN) = 3086
```

Classification Efficiency

Classification Report

```
In [115]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
No	0.87	0.95	0.91	22067
Yes	0.74	0.52	0.61	6372
accuracy			0.85	28439
macro avg	0.80	0.73	0.76	28439
weighted avg	0.84	0.85	0.84	28439

Classification accuracy

```
In [116]: TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]
```

File Edit View Insert Cell Kernel Help

Not Trusted

Python 3 (ipykernel) 









Code

```
In [116]: TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]
```

```
In [117]: # print classification accuracy

classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)

print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))
```

Classification accuracy : 0.8502

Classification error

```
In [118]: # print classification error

classification_error = (FP + FN) / float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))
```

Classification error : 0.1498

Precision

```
In [119]: # print precision score
```

The screenshot displays a Jupyter Notebook interface with the title 'logistic-regression-classifier-tutorial'. The top bar indicates the last checkpoint was 8 minutes ago and the notebook is autosaved. The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Help) and a toolbar with icons for file operations, running cells, and output viewing. The notebook content is divided into sections for different performance metrics:

- Precision**: A code cell (In [119]) calculates precision as $TP / (TP + FP)$ and prints the result: 'Precision : 0.9468'.
- Recall**: A code cell (In [120]) calculates recall as $TP / (TP + FN)$ and prints the result: 'Recall or Sensitivity : 0.8713'.
- True Positive Rate**: A text block states 'True Positive Rate is synonymous with Recall.' followed by a code cell (In [121]) that calculates $TP / (TP + FN)$ and prints 'True Positive Rate : 0.8713'.
- False Positive Rate**: A code cell (In [122]) calculates $FP / (FP + TN)$ and prints 'False Positive Rate : 0.2634'.
- Specificity**: A code cell (In [123]) calculates $TN / (TN + FP)$ and prints 'Specificity : 0.7366'.

The bottom of the image shows the continuation of the notebook interface, including the same menu and toolbar, and the start of the 'Specificity' section.

DISCUSSION and CONCLUSION

The logistic regression model has been applied and executed successfully on the classification problem over the weather dataset of Australia from Kaggle.