

B.TECH. (2020-24)
Artificial Intelligence

Open Ended
LAB File
on
FUNDAMENTALS OF MACHINE LEARNING
[CSE313]



Submitted To
Dr Monika Arora

Submitted By

HITESH	A023119820027	5AI 1
GAURI TYAGI	A023119820028	5AI 1
KUSHAGRA DUBEY	A023119820029	5AI 1
BHOOMIKA SHARMA	A023119820030	5AI 1
SUNIDHI SINGH	A023119820032	5AI 1

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
AMITY SCHOOL OF ENGINEERING AND TECHNOLOGY
AMITY UNIVERSITY UTTAR PRADESH
NOIDA (U.P)

OPEN ENDED EXPERIMENT

Aim

To implement k means clustering algorithm over a dataset

Software Used

Google Colab

Theory

K-Means Clustering is an Unsupervised Learning algorithm, which groups the unlabeled dataset into different clusters. Here K defines the number of pre-defined clusters that need to be created in the process, as if K=2, there will be two clusters, and for K=3, there will be three clusters, and so on. It is a centroid-based algorithm, where each cluster is associated with a centroid. The main aim of this algorithm is to minimize the sum of distances between the data point and their corresponding clusters. The performance of the K-means clustering algorithm depends upon highly efficient clusters that it forms. But choosing the optimal number of clusters is a big task.

Elbow Method

The Elbow method is one of the most popular ways to find the optimal number of clusters. This method uses the concept of WCSS value. **WCSS** stands for **Within Cluster Sum of Squares**, which defines the total variations within a cluster.

To find the optimal value of clusters, the elbow method follows the below steps:

- It executes the K-means clustering on a given dataset for different K values (ranges from 1-10).
- For each value of K, calculates the WCSS value.
- Plots a curve between calculated WCSS values and the number of clusters K.
- The sharp point of bend or a point of the plot looks like an arm, then that point is considered as the best value of K.

Program Code and Output

```
import numpy as np # Linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # for data visualization
import seaborn as sns # for statistical data visualization
%matplotlib inline
```

```
data = 'Live.csv'
```

```
df = pd.read_csv(data)
```

```
df.shape
```

```
(7050, 16)
```

```
df.head()
```

status_type	status_published	num_reactions	num_comments	num_shares	num_likes	num_loves	num_wows	num_hahas	num_sads	num_angrys	Column1
video	4/22/2018 6:00	529	512	262	432	92	3	1	1	0	NaN
photo	4/21/2018 22:45	150	0	0	150	0	0	0	0	0	NaN
video	4/21/2018 6:17	227	236	57	204	21	1	1	0	0	NaN
photo	4/21/2018 2:29	111	0	0	111	0	0	0	0	0	NaN
photo	4/18/2018 3:22	213	0	0	204	9	0	0	0	0	NaN

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7050 entries, 0 to 7049
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   status_id       7050 non-null   object
1   status_type     7050 non-null   object
2   status_published 7050 non-null   object
3   num_reactions   7050 non-null   int64
4   num_comments    7050 non-null   int64
5   num_shares      7050 non-null   int64
6   num_likes       7050 non-null   int64
7   num_loves       7050 non-null   int64
8   num_wows        7050 non-null   int64
9   num_hahas       7050 non-null   int64
10  num_sads        7050 non-null   int64
11  num_angrys      7050 non-null   int64
12  Column1         0 non-null      float64
13  Column2         0 non-null      float64
14  Column3         0 non-null      float64
15  Column4         0 non-null      float64
dtypes: float64(4), int64(9), object(3)
memory usage: 881.4+ KB
```

```
df.isnull().sum()
```

```
status_id      0
status_type    0
status_published 0
num_reactions  0
num_comments    0
num_shares     0
num_likes      0
num_loves      0
num_wows       0
num_hahas      0
num_sads       0
num_angrys     0
Column1       7050
Column2       7050
Column3       7050
Column4       7050
dtype: int64
```

```
df.drop(['Column1', 'Column2', 'Column3', 'Column4'], axis=1, inplace=True)
```

```
df.describe()
```

	num_reactions	num_comments	num_shares	num_likes	num_loves	num_wows	num_hahas	num_sads	num_angrys
count	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000
mean	230.117163	224.356028	40.022553	215.043121	12.728652	1.289362	0.696454	0.243688	0.113191
std	462.625309	889.636820	131.599965	449.472357	39.972930	8.719650	3.957183	1.597156	0.726812
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	17.000000	0.000000	0.000000	17.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	59.500000	4.000000	0.000000	58.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	219.000000	23.000000	4.000000	184.750000	3.000000	0.000000	0.000000	0.000000	0.000000
max	4710.000000	20990.000000	3424.000000	4710.000000	657.000000	278.000000	157.000000	51.000000	31.000000

```
df.corr()
```

	status_type	num_reactions	num_comments	num_shares	num_likes	num_loves	num_wows	num_hahas	num_sads	num_angrys
status_type	1.000000	0.102860	0.320975	0.390910	0.067423	0.388612	0.093844	0.177903	0.081233	0.130989
num_reactions	0.102860	1.000000	0.150843	0.250723	0.994923	0.305003	0.267752	0.176028	0.075138	0.124326
num_comments	0.320975	0.150843	1.000000	0.640637	0.101687	0.521223	0.162394	0.325048	0.236453	0.225184
num_shares	0.390910	0.250723	0.640637	1.000000	0.172492	0.820000	0.407628	0.399826	0.199970	0.312513
num_likes	0.067423	0.994923	0.101687	0.172492	1.000000	0.209308	0.207800	0.120784	0.052169	0.087431
num_loves	0.388612	0.305003	0.521223	0.820000	0.209308	1.000000	0.508798	0.507830	0.207600	0.371001
num_wows	0.093844	0.267752	0.162394	0.407628	0.207800	0.508798	1.000000	0.287756	0.086503	0.183087
num_hahas	0.177903	0.176028	0.325048	0.399826	0.120784	0.507830	0.287756	1.000000	0.141421	0.211910
num_sads	0.081233	0.075138	0.236453	0.199970	0.052169	0.207600	0.086503	0.141421	1.000000	0.142072
num_angrys	0.130989	0.124326	0.225184	0.312513	0.087431	0.371001	0.183087	0.211910	0.142072	1.000000

```
df['status_id'].unique()
```

```
array(['246675545449582_1649696485147474',  
      '246675545449582_1649426988507757',  
      '246675545449582_1648730588577397', ...,  
      '1050855161656896_1060126464063099',  
      '1050855161656896_1058663487542730',  
      '1050855161656896_1050858841656528'], dtype=object)
```

```
len(df['status_id'].unique())
```

```
6997
```

```
df['status_published'].unique()
```

```
array(['4/22/2018 6:00', '4/21/2018 22:45', '4/21/2018 6:17', ...,  
      '9/21/2016 23:03', '9/20/2016 0:43', '9/10/2016 10:30'],  
      dtype=object)
```

```
len(df['status_published'].unique())
```

```
6913
```

```
df.drop(['status_id', 'status_published'], axis=1, inplace=True)
```

```
df.head()
```

	status_type	num_reactions	num_comments	num_shares	num_likes	num_loves	num_wows	num_hahas	num_sads	num_angrys
0	video	529	512	262	432	92	3	1	1	0
1	photo	150	0	0	150	0	0	0	0	0
2	video	227	236	57	204	21	1	1	0	0
3	photo	111	0	0	111	0	0	0	0	0
4	photo	213	0	0	204	9	0	0	0	0

```
X = df
```

```
y = df['status_type']
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
X['status_type'] = le.fit_transform(X['status_type'])
```

```
y = le.transform(y)
```

```
cols = X.columns
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
ms = MinMaxScaler()
```

```
X = ms.fit_transform(X)
```

```
X = pd.DataFrame(X, columns=cols)
```

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=2, random_state=0)
```

```
kmeans.fit(X)
```

```
KMeans(n_clusters=2, random_state=0)
```

```
kmeans.cluster_centers_
```

```
array([[3.28506857e-01, 3.90710874e-02, 7.54854864e-04, 7.53667113e-04,  
       3.85438884e-02, 2.17448568e-03, 2.43721364e-03, 1.20039760e-03,  
       2.75348016e-03, 1.45313276e-03],  
       [9.54921576e-01, 6.46330441e-02, 2.67028654e-02, 2.93171709e-02,  
       5.71231462e-02, 4.71007076e-02, 8.18581889e-03, 9.65207685e-03,  
       8.04219428e-03, 7.19501847e-03]])
```

```
kmeans.inertia_
```

```
237.7572640441955
```

```
labels = kmeans.labels_
```

```
# check how many of the samples were correctly labeled  
correct_labels = sum(y == labels)
```

```
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
```

```
Result: 63 out of 7050 samples were correctly labeled.
```


```
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))
```

```
Accuracy score: 0.01
```

```

from sklearn.cluster import KMeans
cs = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)
    kmeans.fit(X)
    cs.append(kmeans.inertia_)
plt.plot(range(1, 11), cs)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('CS')
plt.show()

```



```

from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=2, random_state=0)

kmeans.fit(X)

labels = kmeans.labels_

# check how many of the samples were correctly labeled
correct_labels = sum(y == labels)

print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))

print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))

```

Result: 63 out of 7050 samples were correctly labeled.
Accuracy score: 0.01

```

kmeans = KMeans(n_clusters=3, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))

```

Result: 138 out of 7050 samples were correctly labeled.
Accuracy score: 0.02

```

kmeans = KMeans(n_clusters=4, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))

```

Result: 4340 out of 7050 samples were correctly labeled.
Accuracy score: 0.62

```

kmeans = KMeans(n_clusters=5, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))

```

Result: 82 out of 7050 samples were correctly labeled.
Accuracy score: 0.01

Discussion and Conclusion

The lesser the model inertia, the better the model fit. So, we use the elbow method to find optimal number of clusters. There is a kink at $k=2$. Hence $k=2$ can be considered a good number of the cluster to cluster this data. So, we have changed the value of k and found relatively higher classification accuracy of 62% with $k=4$. Hence, we can conclude that $k=4$ being the optimal number of clusters.

CRITERIA	TOTAL MARKS	MARKS OBTAINED	COMMENTS
Concept (A)	2		
Implementation (B)	2		
Performance (C)	2		
Total	6		